



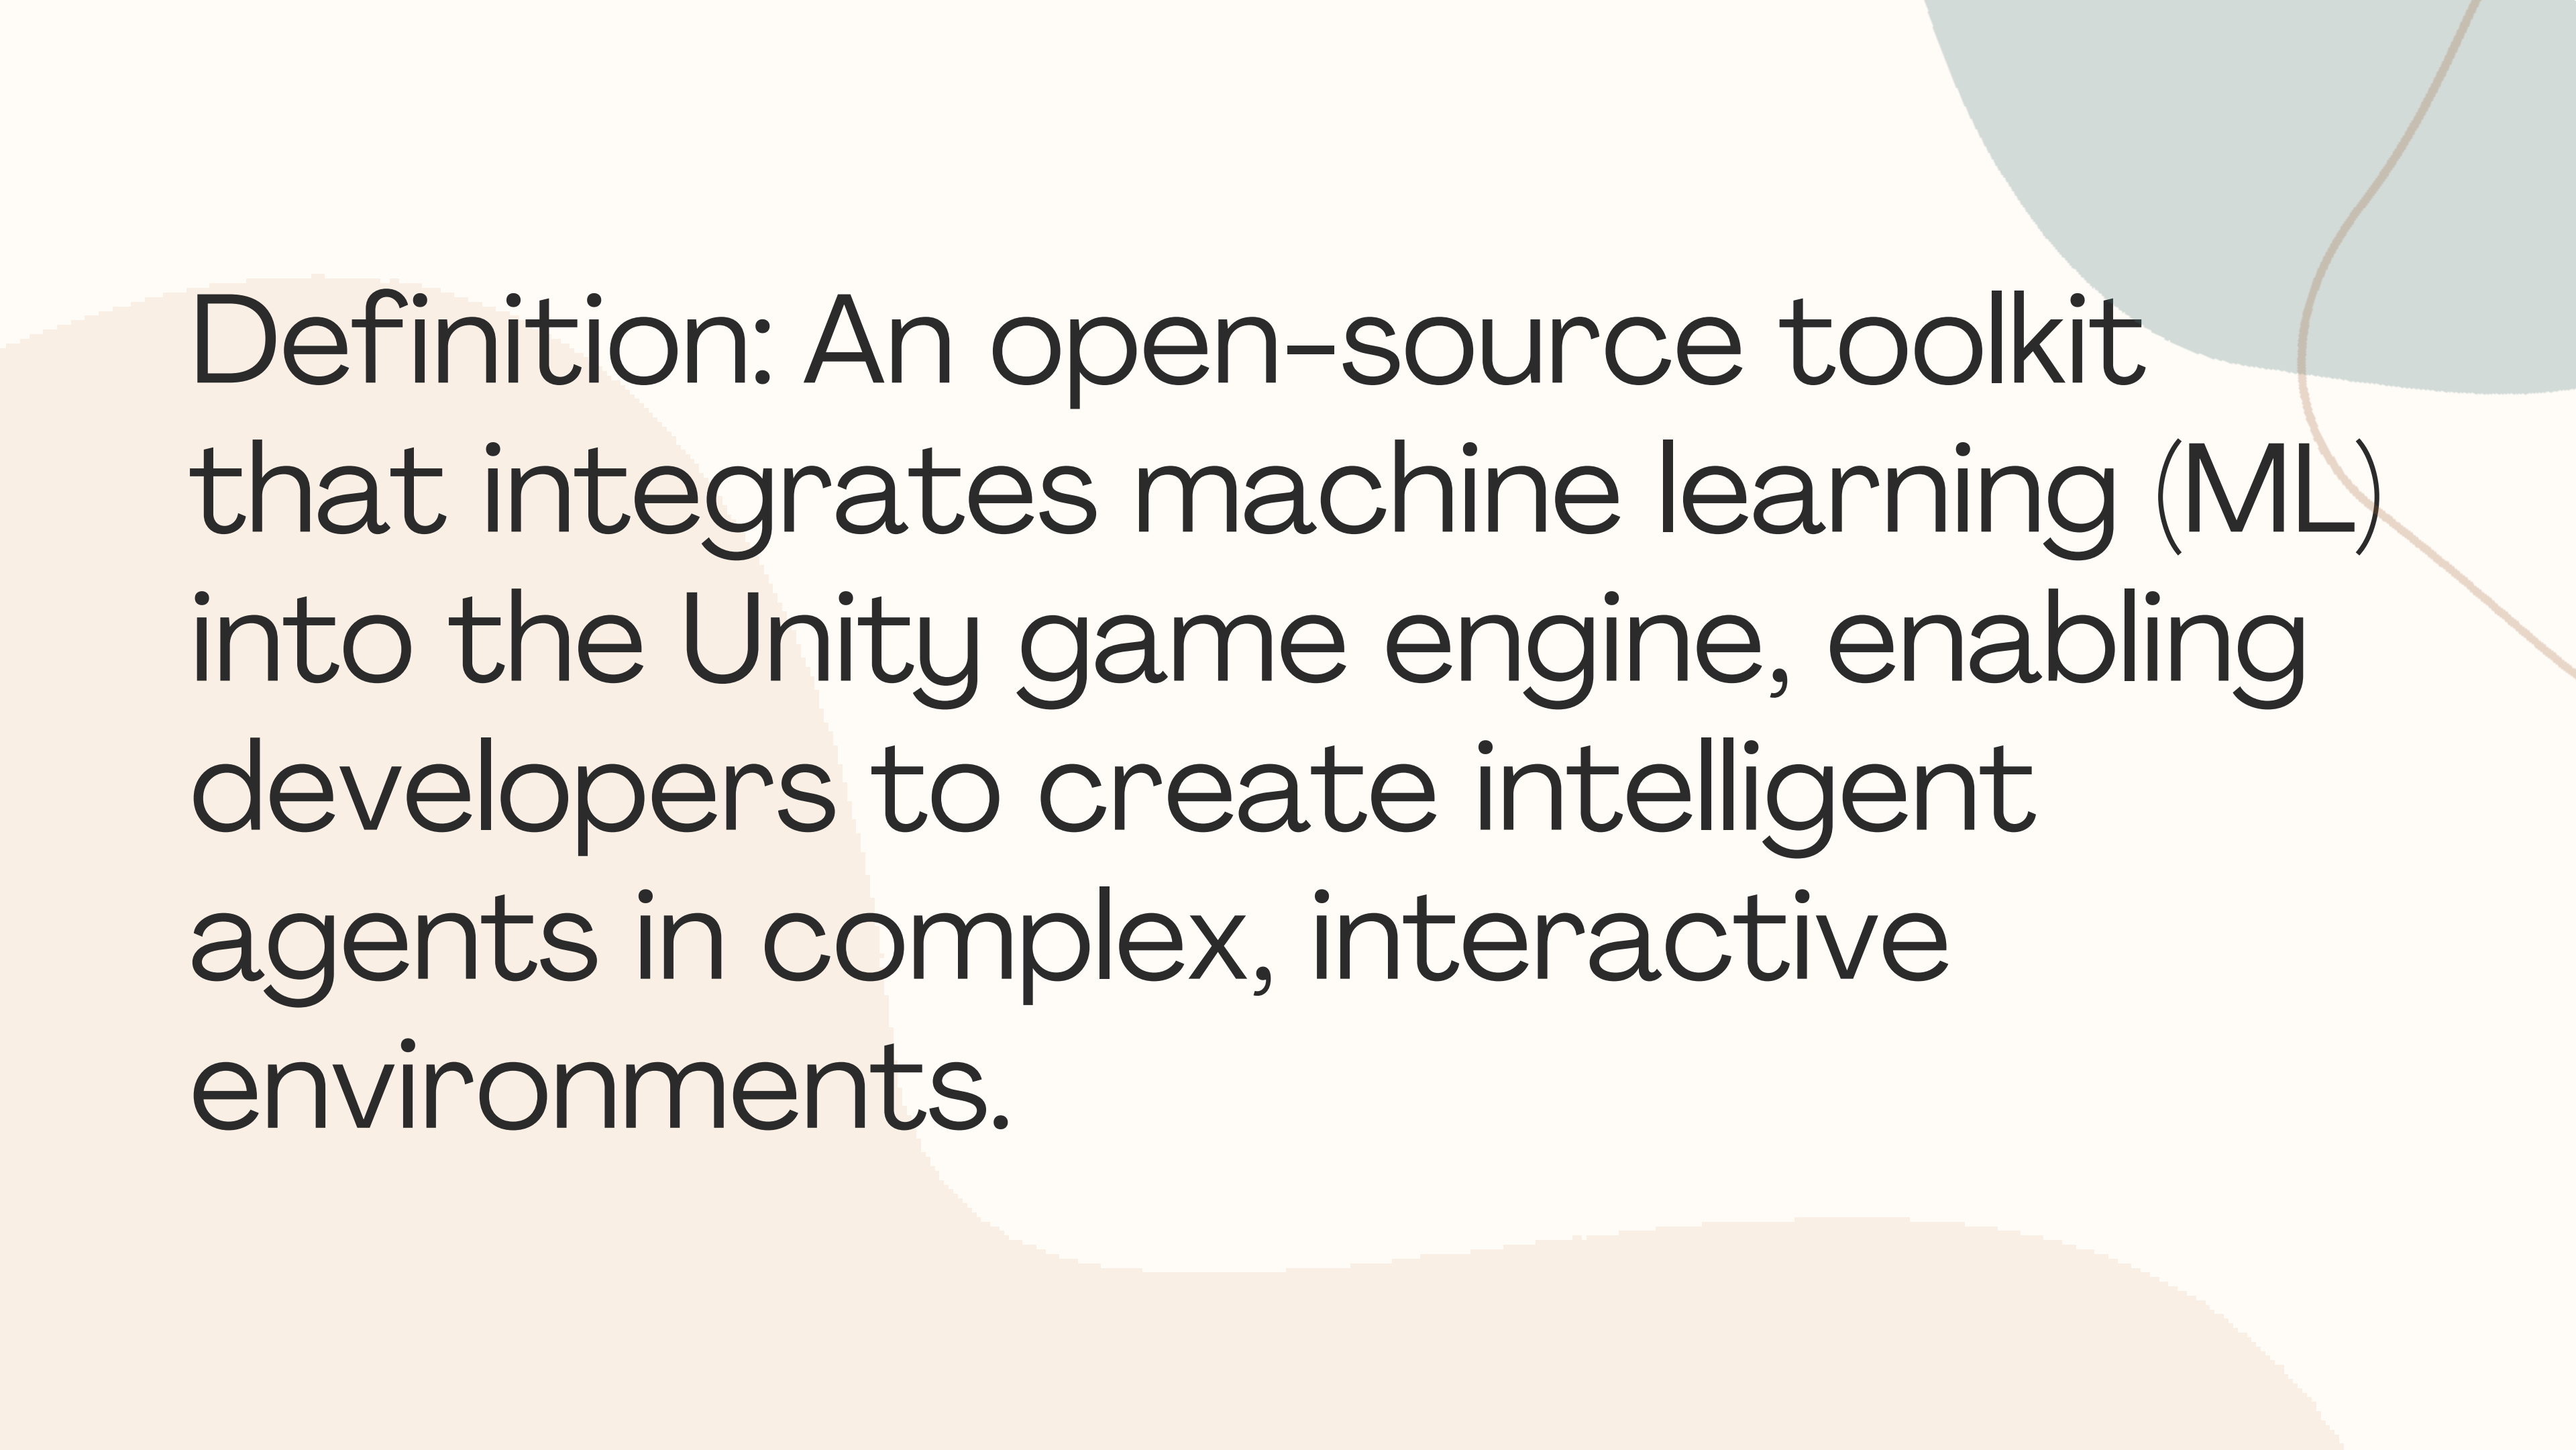
# Week 2

# Theoretical Work

Aim: Understanding fundamentals of Unity ML Agents Reinforcement Learning (RL)



What is Unity ML-Agents?



Definition: An open-source toolkit that integrates machine learning (ML) into the Unity game engine, enabling developers to create intelligent agents in complex, interactive environments.

# Purpose and Goals of Unity ML-Agents

## 1. AI Research and Development:

- Provides a flexible platform for AI research, especially in reinforcement learning, by allowing the creation and evaluation of custom environments.

## 2. Enhance Game Development:

- Integrates advanced AI techniques, creating dynamic and intelligent behaviors for non-player characters (NPCs) to enhance gameplay.

## 3. Support Education and Learning:

- Facilitates hands-on learning of AI and ML concepts, making them more accessible through interactive environments.

## 4. Facilitate Realistic Simulations:

- Develops simulations requiring sophisticated AI, leveraging Unity's 3D rendering and physics engine for realism.

## 5. Promote Innovation in AI Applications:

- Fosters new AI applications across various industries, such as robotics, healthcare, and finance.

## Key Use Cases for ML-Agents

### 1. Training Intelligent NPCs:

- NPCs learn complex behaviors like strategy planning and adapting to player actions.

*Example:* In a stealth game, NPCs are trained to patrol, detect, and strategize.

### 2. Creating Adaptive Gameplay:

- Games adjust difficulty in real-time based on player actions for a personalized experience.

*Example:* AI cars in a racing game adapt to the player's skill level.

### 3. Developing Multi-Agent Systems:

- Supports cooperation and competition in multi-agent environments for richer interactions.

*Example:* AI factions in a strategy game compete for resources or form alliances.

#### 4. Simulating Real-World Scenarios:

- Simulations for research, training, or testing AI decision-making in fields like robotics or autonomous driving.

*Example:* Training autonomous vehicles in virtual environments.

#### 5. Prototyping and Experimentation:

- Allows quick testing and iteration of new AI concepts and algorithms.

*Example:* Prototyping reinforcement learning algorithms in a virtual warehouse.


#### 6. Educational Tools and Interactive Learning:

- Creates interactive learning experiences for teaching AI concepts.

*Example:* Educational games that teach reinforcement learning through guided puzzles.



# Core Components of the ML-Agents Architecture



Overview: The ML-Agents architecture facilitates intelligent agent development within Unity by integrating several key components: Unity Environment, Agent, Academy, Policy, and Python API.



## Key Components and Their Roles

### 1. Unity Environment:

- A virtual world where agents operate, containing all elements (terrains, objects, physics).
- Provides context for agent interactions, sending observations to agents and receiving actions.

### 2. Agent:

- An entity that interacts with the environment, senses it through observations, makes decisions, and takes actions.
- Guided by a Policy that maps observations to actions, learning over time to optimize performance through rewards.

### 3. Academy:

- Manages the training process, controlling environment parameters (like time scale) and managing multiple agents.
- Collects data and communicates with the Python API for processing and updates.

### 4. Policy:

- A strategy for decision-making, typically represented by a neural network.
- Continuously updated during training to maximize cumulative rewards.

### 5. Python API:

- Acts as a bridge between Unity and ML frameworks (TensorFlow, PyTorch).
- Manages communication during training, handles experiment management, logging, and performance visualization.

# Interaction During Training Process

## 1. Initialization:

- Unity environment, Academy, and agents are initialized.
- The Academy configures environment parameters like agent count, time scale, and simulation steps.

## 2. Observation and Action Loop:

- Agents observe their environment, sending data to the Python API.
- The Policy processes observations to decide on actions.

## 3. Action Decision:

- The Policy outputs actions based on input observations.
- Actions are communicated back to the Unity environment.

## 4. Environment Response:

- The environment updates its state based on the action taken and calculates a reward, which is sent to the agent.

## 5. Learning and Policy Update:

- The Python API collects data (observations, actions, rewards) to update the Policy using ML algorithms.
- The Policy is optimized to maximize cumulative rewards over time.

## 6. Repeat:

- The loop repeats until the agent reaches satisfactory performance.

## Interaction During Inference Process

### 1. Fixed Policy Execution:

- The agent uses a fixed, trained Policy to interact with the environment without further learning.
- Decisions are based on the Policy learned during training.

### 2. Real-Time Interaction:

- Actions are executed in real time, and performance is evaluated based on predefined criteria.

### 3. Environment Feedback:

- The agent continues to observe and act based on its fixed Policy until the simulation ends.



# Reinforcement Learning in ML-Agents

# How Reinforcement Learning is Applied in ML-Agents

## Agent-Environment Interaction:

- Observations: The agent perceives the environment by collecting data such as:
  - Vector data: Position, velocity.
  - Visual data: Camera views.
- Action Space: Actions are chosen based on observations:
  - Discrete Actions: Predefined choices (e.g., move left, right, jump).
  - Continuous Actions: Select from a continuous range (e.g., steering angle).
- Rewards: Feedback provided after each action:
  - Incentivizes good behaviors (e.g., reaching a goal).
  - Penalizes undesirable behaviors (e.g., falling off a platform).

# RL Algorithms in ML-Agents

1. Proximal Policy Optimization (PPO):
  - An on-policy algorithm that uses policy gradients to update policies.
  - Ensures stability by controlling the size of policy updates.
  - Default algorithm in ML-Agents due to its balance between stability and performance.
2. Soft Actor-Critic (SAC):
  - An off-policy algorithm that balances expected reward maximization and entropy.
  - Encourages exploration and is effective in continuous action spaces.

These algorithms help agents learn optimal policies to maximize cumulative rewards.

# Training Process in ML-Agents

## 1. Exploration:

- Agents take random or exploratory actions to gather initial data.

## 2. Experience Collection:

- Agents collect experiences (state, action, reward, next state) and store them in a buffer.

## 3. Policy Update:

- After gathering sufficient experiences, agents update their policy.
- Example: PPO minimizes a loss function based on the advantage (difference between expected and actual rewards).

## 4. Iterative Improvement:

- Repeats experience collection and policy update steps.
- Gradually improves agent performance over time.



## Inference Process

- Inference:
  - Uses the trained policy for decision-making in real-time within the Unity environment.
  - The agent selects actions based on the learned policy without further training.





# Proximal Policy Optimization (PPO) and Soft Actor-Critic (SAC) in ML-Agents

# What is Proximal Policy Optimization (PPO)?

## Overview:

- PPO is a policy gradient method designed to improve the stability and efficiency of reinforcement learning.
- It is an on-policy algorithm, meaning it learns from the current policy being executed by the agent.
- PPO ensures that policy updates are not too large, maintaining stability and preventing performance degradation.

## How PPO Works:

- Objective Function: Introduces a clipped objective function to prevent drastic changes in policy.
- Surrogate Objective: Maximizes a surrogate function with a clipping term, ensuring conservative policy updates.
- Advantages: Easy to implement and tune, offering a good balance between sample efficiency and stability.

# What is Soft Actor-Critic (SAC)?

## Overview:

- SAC is an off-policy, actor-critic algorithm that combines value-based and policy-based methods.
- Optimizes the policy while encouraging exploration through an entropy term.
- Well-suited for continuous action spaces and tasks where exploration is crucial.

## How SAC Works:

- Actor-Critic Framework: Uses an actor network (policy) and two critic networks (Q-values) to reduce overestimation bias.
- Entropy Maximization: Includes an entropy term in the objective function to maintain exploration.
- Off-Policy Learning: Learns from past experiences stored in a replay buffer, enhancing sample efficiency.

## Comparison of PPO and SAC

### 1. Stability and Convergence:

- PPO: Ensures stable convergence by using conservative policy updates with a clipping mechanism.
- SAC: Maintains stability through entropy maximization, which is effective in environments with sparse rewards.

### 2. Sample Efficiency:

- PPO: Less sample-efficient as it requires new data for each policy update (on-policy).
- SAC: More sample-efficient due to off-policy learning, allowing data reuse from the replay buffer.

### 3. Exploration:

- PPO: Focuses on refining the current policy; does not explicitly encourage exploration.
- SAC: Explicitly promotes exploration through the entropy term, suitable for environments needing extensive exploration.

#### 4. Complexity and Implementation:

- PPO: Straightforward to implement, with fewer hyperparameters.
- SAC: More complex; involves multiple networks and requires careful tuning, especially for the entropy coefficient.

#### 5. Suitable Applications:

- PPO: Best for tasks requiring stability and simplicity, such as platformer games and navigation tasks with dense rewards.
- SAC: Ideal for continuous action spaces, robotic control, and environments with sparse or deceptive rewards.



# Exploration Strategies in Sparse Reward Environments in Unity ML-Agents

# Introduction to Sparse Reward Environments

Definition:

- Sparse Reward Environments: Scenarios where agents receive feedback or rewards infrequently, making it challenging to learn optimal strategies.
- Challenge: Effective exploration is crucial as agents have limited feedback to guide learning.



## Techniques Used to Encourage Exploration

### 1. Curiosity-Driven Exploration:

- Concept: Agents are intrinsically motivated to explore novel or unpredictable states.
- Mechanism: Assigns “intrinsic rewards” based on the agent’s prediction error, encouraging the exploration of new experiences when external rewards are scarce.

### 2. Entropy Regularization:

- Concept: Promotes randomness in action selection.
- Mechanism: Increases entropy in the agent’s policy, allowing it to explore diverse strategies before settling on an optimal one.
- Benefit: Prevents premature convergence to suboptimal strategies, crucial in complex environments.

### 3. Random Action Selection (Epsilon-Greedy):

- Concept: Mixes exploration with exploitation.
- Mechanism: The agent mostly selects actions predicted to be optimal but occasionally takes random actions to explore new possibilities.
- Benefit: Ensures that the agent does not get stuck in local optima by periodically trying alternative actions.



# Importance of Exploration in Sparse Reward Environments

Why Exploration Matters:

- Discovering Optimal Actions: Enables agents to find the best strategies by exploring various possibilities.
- Avoiding Local Optima: Prevents agents from settling on suboptimal strategies due to limited feedback.
- Learning the Environment: Helps agents gather sufficient data to understand the overall structure of the environment.
- Enhancing Learning Efficiency: Leads to more robust learning outcomes, especially in environments with rare or delayed rewards.



# The Role of the Academy in Unity ML-Agents

# Overview of the Academy

## Definition:

- Academy: A central component in Unity ML-Agents that manages and coordinates the training of agents within a simulated environment.

## Core Functions:

- Simulation Environment Control: Acts as a bridge between the Unity environment and the ML-Agents toolkit, ensuring the environment is correctly managed and agents are synchronized.
- Training Coordination: Manages the timing of episodes, resets the environment, and adjusts its parameters to promote effective learning.

# How the Academy Controls the Simulation Environment

## 1. Episode Management:

- Function: Controls the start and end of training episodes.
- Purpose: Ensures simultaneous resetting of all agents to maintain consistent training.

## 2. Global Parameters Management:

- Function: Manages parameters like the time scale of the simulation.
- Purpose: Allows adjustments to the speed of training while maintaining accuracy in time-dependent elements.

## 3. Communication with Python:

- Function: Facilitates communication between the Unity environment and the Python training process.
- Purpose: Sends observations to the Python API and receives actions in return to update agents' behavior based on the latest policy.

# Configuring the Academy for Optimal Training

## 1. Time Scale Adjustment:

- Parameter: timeScale
- Use: Adjusts the speed of the simulation (e.g., faster training with a higher time scale).
- Consideration: Ensure accurate physics and time-dependent simulation elements.

## 2. Max Steps Configuration:

- Parameter: Max Steps
- Use: Sets the maximum number of steps per episode.
- Goal: Balance learning efficiency and resource usage by adjusting episode length.

## 3. Randomization and Curriculum Learning:

- Purpose: Introduces randomness or gradually increases task difficulty.
- Benefit: Encourages generalized learning and prevents overfitting.

## 4. Reset Parameters:

- Function: Resets specific parts of the environment or introduces random variations.
- Goal: Creates varied training experiences to promote robust learning.



# Unity ML-Agents: Supporting Multi-Agent Environments

# Multi-Agent Scenarios in Unity ML-Agents

- Overview: Unity ML-Agents supports various multi-agent scenarios where multiple agents interact within the same environment.

## Types of Scenarios:

1. Cooperative Scenarios: Agents collaborate to achieve a shared goal.
  - Example: Robots carrying an object together.
2. Competitive Scenarios: Agents compete against each other for rewards or resources.
  - Example: Racing game where agents race to reach the finish line first.
3. Mixed Scenarios: Some agents cooperate while others compete.
  - Example: Team-based sports games where agents cooperate within teams and compete against other teams.



## Multi-Agent Support Features in Unity ML-Agents

- Separate Agents:
  - Each agent has its own policy and learns independently.
  - Suitable for competitive or independent learning scenarios.
- Shared Policies:
  - Multiple agents share a common policy.
  - Reduces complexity by using a single policy for agents that should behave similarly.
- Custom Reward Functions:
  - Each agent can have its own reward function.
  - Allows nuanced learning tailored to individual agent objectives.
- Communication:
  - Agents may share observations or communicate.
  - Enhances coordination in cooperative tasks.



# Training in Multi-Agent Environments

- Synchronous Training:
  - The Academy ensures that all agents interact with the environment simultaneously.
  - Enables synchronized training across multiple agents.
- Policies:
  - Agents can have individual or shared policies.
  - Multiple Policies: Adds complexity but allows for specialized behavior.
- Rewards:
  - Can be agent-specific or global.
  - Depends on whether agents are collaborating or competing.

# Challenges in Multi-Agent Reinforcement Learning

## 1. Non-Stationarity:

- The environment constantly changes as all agents learn and adapt.
- From an individual agent's perspective, this makes the environment non-stationary, complicating policy optimization.

## 2. Credit Assignment:

- Difficult to assign appropriate credit (or blame) for group outcomes.
- Complicates reward assignment and can slow learning.

## 3. Coordination:

- Agents must learn to coordinate actions, which is challenging with different goals or incomplete information.

## 4. Competition:

- Agents must balance maximizing their own rewards with minimizing opponents' rewards.
- Introduces adversarial dynamics in learning.

# Techniques for Multi-Agent Reinforcement Learning

1. Centralized Training with Decentralized Execution (CTDE):
  - Agents are trained together in a shared environment.
  - During execution, each agent uses its own policy.
2. Curriculum Learning:
  - Agents start with simpler tasks or fewer agents.
  - Progress to more complex scenarios as their skills improve.



# Designing Custom Reward Functions in Unity ML-Agents

# Importance of Custom Reward Functions

- Definition: Custom reward functions guide agents in learning specific tasks effectively.
- Goal: Ensure agents learn the desired behavior efficiently, avoiding unintended outcomes.

# Key Considerations for Designing Reward Functions

## 1. Alignment with Desired Behavior:

- Ensure the reward function reflects the task's true objectives.
- Example: For a maze, reward reaching the exit rather than random movement.

## 2. Avoiding Exploitation:

- Prevent agents from exploiting the reward system.
- Example: If rewarding object collection, specify constraints to avoid agents repeatedly collecting the same object.

## 3. Sparse vs. Dense Rewards:

- Sparse Rewards: Given only for significant accomplishments.
- Dense Rewards: Provide frequent feedback.
- Combine both to balance feedback and maintain focus on the objective.

## 4. Scaling and Consistency:

- Scale rewards appropriately to avoid overwhelming or under-rewarding the agent.
- Ensure that rewards do not overshadow or ignore important behaviors.

# Reward Shaping Examples

## 1. Maze Navigation:

- Basic Reward: Reward for reaching the exit.
- Reward Shaping: Provide incremental rewards for reducing the distance to the exit. Apply small penalties for moving in the wrong direction or staying idle.

## 2. Object Collection:

- Basic Reward: Reward for each object collected.
- Reward Shaping: Higher rewards for collecting objects in a specific sequence or within a time limit. Bonus for completing the collection efficiently.

## 3. Pole Balancing:

- Basic Reward: Reward for keeping the pole balanced.
- Reward Shaping: Continuous rewards for each second the pole remains balanced, with penalties for large or sudden movements.

## 4. Autonomous Driving:

- Basic Reward: Reward for completing a lap or reaching the destination.
- Reward Shaping: Incremental rewards for maintaining correct speed, staying in lanes, and avoiding collisions. Penalties for sharp turns or crashes.



# Integration of Unity ML-Agents with External ML Frameworks



## Overview of Unity ML-Agents Integration

- Unity ML-Agents: A toolkit for training intelligent agents in Unity using reinforcement learning (RL) and imitation learning.
- Integration Purpose: To leverage external ML frameworks like TensorFlow and PyTorch for advanced algorithms and custom models.

## Role of the Python API in Integration

### 1. Communication:

- Utilizes a gRPC-based interface to exchange data between Unity and the ML framework.
- Transfers states, actions, rewards, and agent observations to facilitate training.

### 2. Training Control:

- The Python API defines the training loop and handles data processing.
- Allows integration with TensorFlow or PyTorch for custom models and algorithms.

### 3. Data Handling:

- Collects and preprocesses observation data from Unity.
- Feeds the processed data to the ML framework for training.
- Trained models generate actions that are sent back to Unity to update agent behavior.

### 4. Custom Algorithm Implementation:

- Enables implementation of custom training algorithms not available in Unity ML-Agents.
- Acts as a bridge, giving external frameworks control over the learning process.

## Advantages of Using External ML Frameworks

### 1. Flexibility in Model Design:

- Access to robust libraries and tools for designing complex neural networks.
- Allows experimentation with diverse architectures beyond Unity ML-Agents' defaults.

### 2. Advanced Features and Custom Algorithms:

- Leverages state-of-the-art tools, optimizers, and algorithms (e.g., PPO, A2C).
- Utilizes GPU acceleration, mixed-precision training, and distributed computing.

### 3. Community Support and Resources:

- Access to extensive documentation, pre-trained models, and a large community for support.

### 4. Interoperability:

- Enables cross-platform compatibility and seamless integration with existing ML pipelines.

# Challenges of Using External ML Frameworks

## 1. Performance Overhead:

- Potential latency due to communication between Unity and the external ML framework.
- Overhead can be significant in complex environments or with large models.

## 2. Complexity in Setup and Maintenance:

- Requires understanding both Unity ML-Agents and external ML framework APIs.
- Maintaining compatibility across different software versions and dependencies can be challenging.

## 3. Debugging and Monitoring:

- Debugging involves both the Unity environment and the ML framework.
- May need custom tools for efficient visualization and logging of data.

## 4. Resource Requirements:

- Training complex models may need significant computational resources (e.g., GPU or TPU).
- Ensuring optimal use of resources is crucial, especially in large-scale experiments.



Thanks for listening!